

ABHIRUP DAS

Data Engineer, Azure Data Platform

Databricks — PySpark — SQL — Power BI

Email: abhirupdas99official@gmail.com | Phone: +1 647-535-0832 | Toronto, Ontario, Canada

LinkedIn: linkedin.com/in/abhirupdas99 | GitHub: github.com/abhirupdas99

Professional Summary

- Data Engineer with **5+ years** designing and maintaining **scalable ETL pipelines, dimensional data models, and performance reporting solutions** across financial and operational domains
- Delivered pipelines processing **35000+ weekly financial transactions**, cut Power BI dashboard load time from 18s to 11s through star-schema redesign, and reduced data defects 22% with validation and reconciliation.
- Hands-on with **Azure Databricks, Azure Data Factory, PySpark, SQL, Power BI, and Python**. **Authorized to work anywhere in Canada without sponsorship** and available for full-time Data Engineer roles immediately.

SKILLS

- **Data Processing & Transformation:** Apache Spark (PySpark, Spark SQL),dbt
- **Pipeline Orchestration & Ingestion:** ETL/ELT design, incremental loading, CDC, batch and streaming (Azure Data Factory, Auto Loader, Apache Kafka)
- **Lakehouse & Storage Architecture:** Delta Lake, medallion architecture (Bronze/Silver/Gold), ADLS Gen2 (Databricks),Unity Catalog
- **Data Modeling:** Dimensional modeling, star schema, slowly changing dimensions, normalization
- **Cloud Platforms:**Azure (Databricks, ADF, ADLS Gen2, Synapse),AWS (S3, Glue, Redshift),Snowflake
- **BI & Reporting:** Power BI (semantic models, DAX, datasets, KPI dashboards), Tableau
- **Engineering Practices:** Git, Jira, Agile/Scrum, CI/CD, unit testing
- **Languages :** Python, SQL

WORK EXPERIENCE

Business Intelligence

09/2025 – Present

Storefront Stays Ltd

Windsor, ON, CA

- Defined and tracked 20+ business **KPIs** using standardized Excel models, reducing manual reporting effort by 30% and creating a consistent reporting process across regional operations.
- Authored Power BI semantic model with **15+ DAX measures** (export volume, on-time-shipment %, country-level revenue) consumed by executive dashboards across 6+ markets, replacing 4 manual Excel reports.
- Developed **Power BI** and **Tableau** dashboards on top of the semantic model, giving stakeholders same-day visibility into export orders, shipment delays, and regional performance, enabling faster weekly operational reviews across 6+ markets.

Data Engineer

10/2024 – 08/2025

Windsor Analytics Ltd

Windsor, ON, CA

- Designed and built ETL pipelines in **Azure Databricks** using **PySpark** to process 35,000+ weekly financial transactions, handling ingestion, transformation, validation, and aggregation before publishing curated datasets that reduced dashboard latency by 25%.
- Built streaming ingestion pipelines using **Apache Kafka**, processing real-time business events before landing them in Delta Lake for downstream transformations and reporting.
- Implemented **Delta Lake** data quality framework, including schema enforcement, constraint validation, alerting, and complete lineage tracking via **Unity Catalog** that maintained 100% schema compliance across the Gold layer.
- Integrated data from **RESTful APIs** into **Snowflake**, built dimensional models, and published curated datasets for Power BI reporting using reusable transformation models.
- Improved stakeholder alignment by translating complex technical data models and ETL logic into clear business-facing documentation and walkthrough sessions.

Data Engineer

06/2023 – 08/2024

Accenture Solutions Pvt Ltd

India

- Validated telecom customer and order data flowing from Salesforce CRM to warehouse tables using **SQL** to verify mappings, joins and transformation rules across Accounts, orders and product objects.
- Performed end-to-end reconciliation across **Salesforce, AWS S3, AWS Glue, and warehouse tables**, validating record counts, key relationships, null handling, and order lifecycle scenarios to maintain data accuracy.
- Loaded transformed datasets from AWS Glue into **Amazon Redshift**, supporting analytical workloads and improving query performance through dimensional data models.
- Investigated pipeline failures using **Apache Airflow** and integration logs, identifying schema changes, delayed extracts, and partial data loads before coordinating successful reprocessing.

- Compared warehouse data with Salesforce transactions and API responses to verify billing, reporting, and customer service status before production releases.
- Supported release cycles with sanity, regression and production data verification, improving overall data reliability across CRM and reporting systems.
- Modeled slowly changing customer and product dimensions using **SCD Type 2 in SQL Server**, capturing effective-from/effective-to dates across 120K records and enabling 14 months of point-in-time reporting for finance.

Junior Software Engineer

Accenture Solutions Pvt Ltd

07/2021 – 05/2023

India

- Designed scalable SQL data architectures, built **star-schema** dimensional models, and developed **dbt transformation models** to standardize business rules and improve consistency across downstream analytical datasets.
- Wrote **PyTest** unit tests reaching 80% coverage on transformation logic and integrated test runs into **CI/CD pipeline (Azure DevOps)**, catching schema regressions before production deploy across 3+ pipelines.
- Applied **Python** root-cause analysis (regression, time-series forecasting) to a production scheduling anomaly, delivered findings to operations leadership that drove a **20% throughput improvement**.
- Deployed modular, production-grade pipelines using **CI/CD workflows (GitHub Actions)**, automated testing and version-controlled modular notebooks that enabled zero-downtime releases and full auditability across deployments.

PROJECT EXPERIENCE

Enterprise Azure Data Platform Integrating 8 Data Sources with Real-Time Streaming, CDC & Medallion Architecture

01/2026 – 04/2026

Developer, Ontario

- Built Azure lakehouse on ADLS Gen2, ingesting 8 sources (REST APIs, SQL Server CDC feeds, Event Hubs streams, CSV/JSON files) via Azure Data Factory and Databricks (PySpark) into Delta Lake tables partitioned by ingestion date for query pruning within a Medallion (Bronze-Silver-Gold) architecture.
- Wrote **PySpark** and **Spark SQL** transformations on Silver layer with Structured Streaming + Auto Loader for incremental ingestion, handling schema drift and late-arriving events with watermarks set to 30 minutes
- Implemented **SCD Type 2** customer and product dimensions in Gold-layer Delta tables using **PySpark MERGE INTO**, preserving 100% of historical attribute changes across 500K dimension records for time-travel reporting
- Wired Azure DevOps CI/CD pipeline deploying Databricks notebooks via **Repos and Azure Synapse serverless SQL** views over Gold tables, cutting Power BI report-to-data latency from 12 hours to 15 minutes
- Configured **Unity Catalog** with table-level lineage, column-level access controls, and audit logs across 40+ tables; integrated KQL alerts in Azure Monitor for SLA breaches and data freshness violations
- Built **Power BI** dashboards on curated Gold tables, giving business users access to operational and financial metrics while enforcing RBAC and row-level security.

Personal Financial Analysis using Power BI

09/2024 – 07/2025

Developer, Ontario

- Implemented interactive Power BI dashboard using twelve months bank statement data, visualizing income, expenses, and savings trends efficiently.
- Tracked monthly income, expenditure, and savings patterns, enabling strategic financial planning and resource allocation for precise budget optimization.

World Layoffs Data Cleaning Project with MySQL

01/2023 – 08/2023

Sole Contributor, Ontario

- Cleaned global layoffs dataset (2020–2023), ensuring accurate analysis of workforce reduction across companies, industries, and countries systematically.
- Interpreted staging tables preserving raw datasets, enabling repeatable and reliable data cleaning workflows, improving overall data integrity and accuracy.

EDUCATION

Master of Applied Computing

University of Windsor, Ontario

09/2024 – 12/2025

Bachelor of Technology (B.Tech)

Narula Institute of Technology, India

08/2017 – 06/2021

Courses & Certifications

- **Microsoft Certified: Azure Data Fundamentals (DP-900)**
- **Microsoft Certified: Azure Data Engineer Associate (DP-203)**
- **Azure Spark Databricks Essential Training** – LinkedIn Learning
- **Learning Hadoop & Learning MongoDB** – LinkedIn Learning
- **Power BI Essential Training** – LinkedIn Learning