

ABHIRUP DAS

Data Engineer, Azure Data Platform

Databricks — PySpark — SQL — Power BI

Email: abhirupdas99official@gmail.com | Phone: +1 647-535-0832 | Toronto, Ontario, Canada

LinkedIn: linkedin.com/in/abhirupdas99 | GitHub: github.com/abhirupdas99

Professional Summary

- Data Engineer with **5+ years** designing and maintaining **scalable ETL pipelines, dimensional data models, and performance reporting solutions** across financial and operational domains
- Delivered pipelines processing **90,000+ weekly financial transactions**, cut Power BI dashboard load time from 18s to 11s through star-schema redesign, and reduced data defects 22% with validation and reconciliation.
- Hands-on with **Azure Databricks, Azure Data Factory, PySpark, SQL, Power BI, and Python**. **Authorized to work anywhere in Canada without sponsorship** and available for full-time Data Engineer roles immediately.

SKILLS

- **Cloud & Data Platforms:** Azure Databricks, Azure Data Factory, ADLS Gen2, Azure Synapse, Azure Key Vault, Azure Monitor, Azure DevOps
- **Data Modeling & Warehousing:** Medallion Architecture, Lakehouse, Dimensional Modeling, Star/Snowflake Schema, SCD Type 1/2/3, Data Governance, Data Lineage
- **Big Data & Streaming:** Apache Spark, PySpark, Spark SQL, Delta Lake, Structured Streaming, Event Hubs
- **Python Libraries:** pandas, NumPy, scikit-learn
- **BI & Reporting:** Power BI (semantic models, DAX, datasets, KPI dashboards), Tableau
- **Engineering Practices:** Git, Jira, Agile/Scrum, CI/CD, unit testing

WORK EXPERIENCE

Data Engineer

06/2023 – 08/2024

Accenture Solutions Pvt Ltd

India

- Designed star schema with 1 fact and 5 conformed dimensions in SQL Server, replacing flat-table queries; cut **Power BI dashboard load time from 18s to 11s** (40% reduction) for 30+ daily users.
- Wired SQL stored procedures to Power BI contract-renewal visuals **tracking 8K+ active contracts**, surfacing at-risk accounts 30 days earlier and contributing to a 10-percentage-point lift in renewal close rate.
- Delivered biweekly Tableau evaluation reports analyzing path recency, asset relevancy, accuracy metrics, assisting data scientists, guiding iterative model deployment improvements.
- Ran regression and time-series forecasts in **Python (pandas, NumPy, scikit-learn)** on production volume data, identifying bottleneck processes that drove a 20% efficiency lift after operations team rerouted 4 production lines.
- Modeled slowly changing customer and product dimensions using **SCD Type 2 in SQL Server**, capturing effective-from/effective-to dates across 120K records and enabling 14 months of point-in-time reporting for finance.

Junior Software Engineer

07/2021 – 05/2023

Accenture Solutions Pvt Ltd

India

- Designed and built **production ETL pipelines on Azure Databricks and Apache Spark** processing 90,000+ weekly financial transactions, implementing **data ingestion, transformation, validation, conforming, and aggregation** across the full data lifecycle with scheduled refresh cycles.
- Wrote PyTest unit tests reaching 80% coverage on transformation logic and integrated test runs into CI/CD pipeline (Azure DevOps), catching schema regressions before production deploy across 10+ pipelines.
- Built and monitored **10+ production data pipelines**, performing issue triage, root-cause analysis of data failures, and implementing controls and remediation steps to maintain **data availability and SLA compliance**.
- Designed **physical database schemas** with stored procedures, data retrieval strategies, and query optimization, delivering performant, well-structured data for downstream reporting and analytics consumers.

Business Intelligence Intern

09/2025 – Present

Storefront Stays Ltd — Internship (Remote)

Windsor, ON, CA

- Built KPI tracking models in Excel to monitor **20+ performance indicators** and reduced reporting time by 25%.
- Authored Power BI semantic model with **15+ DAX measures** (export volume, on-time-shipment %, country-level revenue) consumed by executive dashboards across 15 markets, replacing 4 manual Excel reports.
- Built **Power BI dashboards across 15+ markets** to support weekly reviews, resulting in same-day visibility on export orders and delays.

Data Engineer Intern

10/2024 – 08/2025

Windsor Analytics Ltd

Windsor, ON, CA

- Designed **scalable SQL data architectures**, implemented star-schema dimensional models, defined transformation logic and delivered data integration pipelines that **reduced dashboard latency by 40%** for downstream reporting consumers.
- Developed **ETL workflows with scheduled refreshes and validation controls**, applied transformation, conforming, and reconciliation logic, **reducing reporting defects by 22%** across financial datasets.
- Identified and **automated manual data processes**, engineered SQL-driven early-warning signals and reconciliation pipelines that replaced manual checks, contributing to a **10% improvement in downstream business outcomes**.
- Applied **iterative development across the full data lifecycle**, data profiling, ingestion, transformation, modelling, cleansing, and aggregation, partnering with engineering teams to deliver well-structured data assets for reporting consumers.
- Applied **Python root-cause analysis** (regression, time-series forecasting) to a production scheduling anomaly, delivered findings to operations leadership that drove a **20% throughput improvement**.

PROJECT EXPERIENCE

Enterprise Azure Data Platform Integrating 8 Data Sources with Real-Time Streaming, CDC & Medallion Architecture

01/2026 – 04/2026

Developer, Ontario

- Built Azure lakehouse on ADLS Gen2, ingesting 8 sources (REST APIs, SQL Server CDC feeds, Event Hubs streams, CSV/JSON files) via Azure Data Factory and Databricks (PySpark) into Delta Lake tables partitioned by ingestion date for query pruning within a Medallion (Bronze-Silver-Gold) architecture.
- Wrote **PySpark** and **Spark SQL** transformations on Silver layer with Structured Streaming + Auto Loader for incremental ingestion, handling schema drift and late-arriving events with watermarks set to 30 minutes
- Implemented **SCD Type 2** customer and product dimensions in Gold-layer Delta tables using **PySpark MERGE INTO**, preserving 100% of historical attribute changes across 500K dimension records for time-travel reporting
- Wired Azure DevOps CI/CD pipeline deploying Databricks notebooks via **Repos and Azure Synapse serverless SQL** views over Gold tables, cutting Power BI report-to-data latency from 12 hours to 15 minutes
- Configured Unity Catalog with table-level lineage, column-level access controls, and audit logs across 40+ tables; integrated KQL alerts in Azure Monitor for SLA breaches and data freshness violations
- Delivered business-ready insights through optimized data models and Power BI dashboards, while implementing secure, production-grade practices such as RBAC, Key Vault integration, and row-level security.

Personal Financial Analysis using Power BI

09/2024 – 07/2025

Developer, Ontario

- Implemented interactive Power BI dashboard using twelve months bank statement data, visualizing income, expenses, and savings trends efficiently.
- Tracked monthly income, expenditure, and savings patterns, enabling strategic financial planning and resource allocation for precise budget optimization.

World Layoffs Data Cleaning Project with MySQL

01/2023 – 08/2023

Sole Contributor, Ontario

- Cleaned global layoffs dataset (2020–2023), ensuring accurate analysis of workforce reduction across companies, industries, and countries systematically.
- Interpreted staging tables preserving raw datasets, enabling repeatable and reliable data cleaning workflows, improving overall data integrity and accuracy.

Credit Risk Prediction and Explainability Platform

05/2025 – 08/2025

Sole Contributor, Ontario

- Built credit-default classification pipeline in Python using pandas + NumPy for feature engineering on 30K+ loan records, scikit-learn for preprocessing and train/test split, and XGBoost/LightGBM with SHAP feature attributions achieving 0.84 ROC-AUC
- Evaluated predictive model performance using ROC-AUC, F1-score, and accuracy metrics, optimizing model selection and ensuring robust predictive capability.

EDUCATION

Master of Applied Computing

09/2024 – 12/2025

University of Windsor, Ontario

Bachelor of Technology (B.Tech)

08/2017 – 06/2021

Narula Institute of Technology, India

Courses & Certifications

- **Microsoft Certified: Azure Data Fundamentals (DP-900)**
- **Microsoft Certified: Azure Data Engineer Associate (DP-203)**
- **Azure Spark Databricks Essential Training** – LinkedIn Learning
- **Learning Hadoop & Learning MongoDB** – LinkedIn Learning
- **Power BI Essential Training** – LinkedIn Learning